

Artists in the USA

Jonathan Fung

09/27/22

Contents

Import	2
Impute	3
Diagnostics	5
Utils	8
Explore	9
GIS	12
Resources	16

Import

Source: [tidytuesday - September 27, 2022](#)

```
library(tidyverse)
library(mice)

tuesdata <- tidytuesdayR::tt_load('2022-09-27')
artists <- tuesdata$artists
options(crayon.enabled = FALSE)
```

```
head(artists)
```


A tibble: 6 × 7

	state	race	type	all_workers_n	artists_n	artists_share	location_quotie...
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	Alabama	Hispa...	Arch...	88165	45	0.000510	0.875
2	Alaska	Hispa...	Arch...	26875	15	0.000558	0.957
3	Arizona	Hispa...	Arch...	1033370	270	0.000261	0.448
4	Arkansas	Hispa...	Arch...	101405	NA	NA	NA
5	California	Hispa...	Arch...	7470730	3870	0.000518	0.888
6	Colorado	Hispa...	Arch...	594525	200	0.000336	0.577

Impute

We can inspect the data's missing values. There are 1909 complete records, and 1471 (17 + 755 + 699) incomplete records, rows that contain an NA value.

```
md.pattern(artists, rotate.names = TRUE)
```



Since this dataset has a high amount of missing data, we use Multiple Imputation by Chained Equations (MICE).

```
imp_base <- mice(artists, m = 5, maxit = 0, print = FALSE)
mypred <- imp_base$pred
mypred[c("state", "race", "type"),] <- 0 # constants
## Sect 9: https://www.gerkovink.com/miceVignettes/Passive_Post_processing/Passive_imp_
  → utation_post_processing.html
mypred[c("location_quotient", "artists_n"), "artists_share"] <- 0 # linear combination
mymeth <- imp_base$meth
mymeth["artists_share"] <- "~I(artists_n/all_workers_n)"

imp <- mice(artists, pred = mypred, meth = mymeth, maxit = 35, print = F, seed = 123)
artists_mi <- imp %>% complete() %>% as_tibble()
head(artists_mi)
```

Warning message:

Number of logged events: 3

A tibble: 6 × 7

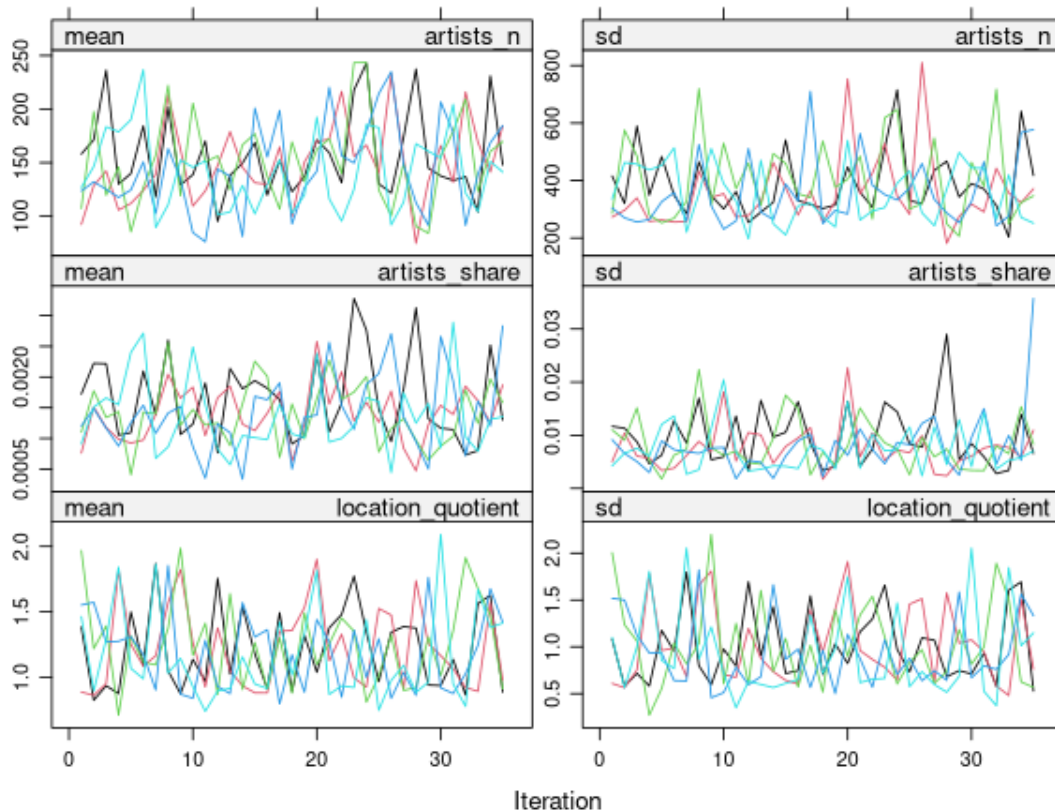
	state	race	type	all_workers_n	artists_n	artists_share	location_quotie...
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	Alabama	Hispa...	Arch...	88165	45	0.000510	0.875
2	Alaska	Hispa...	Arch...	26875	15	0.000558	0.957

3	Arizona	Hispa...	Arch...	1033370	270	0.000261	0.448
4	Arkansas	Hispa...	Arch...	101405	45	0.000444	0.962
5	California	Hispa...	Arch...	7470730	3870	0.000518	0.888
6	Colorado	Hispa...	Arch...	594525	200	0.000336	0.577

Diagnostics

We can check the general convergence of the MICE algorithm. Each row represents a variable that had missing data, with the LHS being the mean, and RHS being the standard deviation. A sufficient diagnostic plot should have no trends across iterations, which these mostly do.

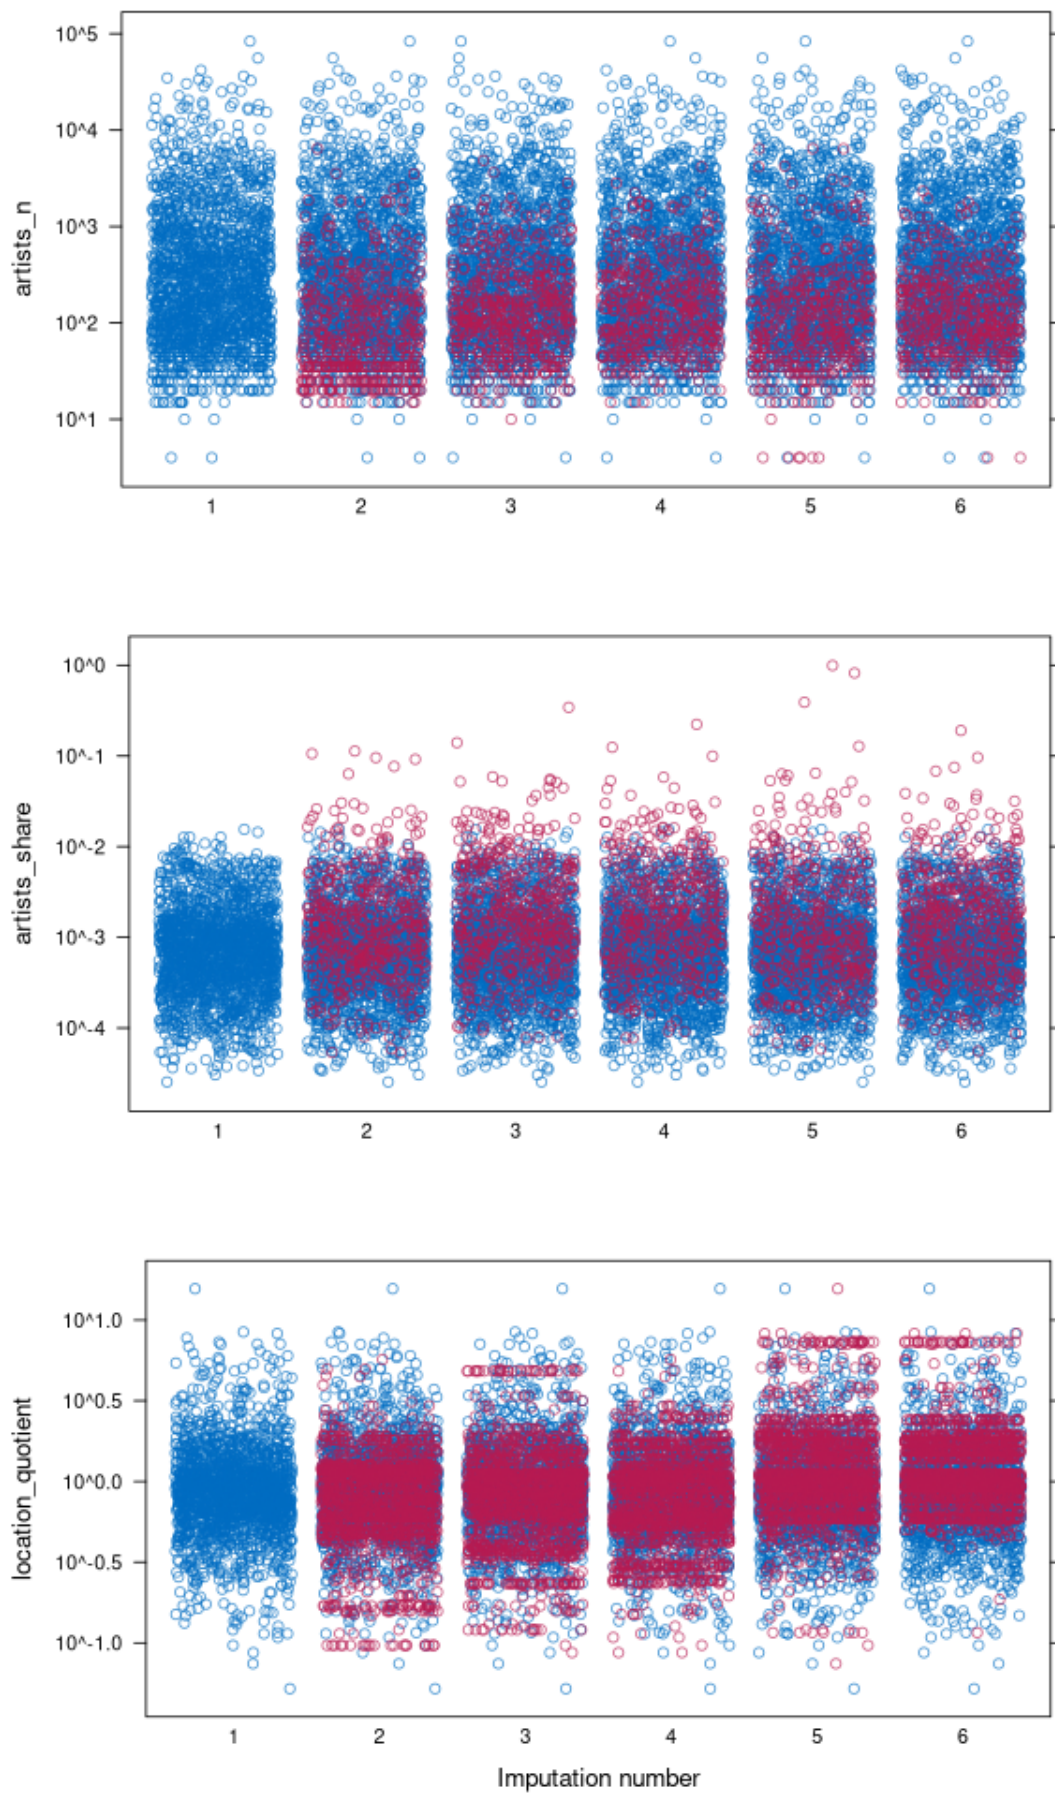
```
plot(imp)
```



Next, we can show that the imputed values (Red) are within range (*plausible*) of the original values (Blue):


```
## https://stackoverflow.com/questions/2540129/lattice-multiple-plots-in-one-window
sp1 <- stripplot(imp, artists_n~.imp, scales=list(y = list(log = 10)), xlab = "",
  ↪ factor = 2)
sp2 <- stripplot(imp, artists_share~.imp, scales=list(y = list(log = 10)), xlab = "",
  ↪ factor = 2)
sp3 <- stripplot(imp, location_quotient~.imp, scales=list(y = list(log = 10)), factor
  ↪ = 2)

#library(gridExtra)
gridExtra::grid.arrange(sp1, sp2, sp3)
```



Similarly, we can inspect their distribution:

```
densityplot(imp, scales=list(x = list(log = 10)), layout = c(3,1))
```



Utils

Group up states into regions: Northeast, South, North Central, West

```
assignRegion <- function(state) {
  lvls <- levels(datasets::state.region)

  res <- ifelse(state == "District of Columbia" || state == "Puerto Rico",
    datasets::state.region[1],
    datasets::state.region[match(state, datasets::state.name)]
  )
  return(factor(lvls[res], levels = lvls))
}

df <- artists_mi %>% mutate(region = Vectorize(assignRegion)(state))
```

Function to estimate the mode(s) of a univariate distribution, particularly multimodal.

```
## from package of the same name
ModEstM <- function(x, ...){
  Density = stats::density(x, ...)
  ###
  data.frame(abciss = Density$x,
    density = Density$y) |>
  dplyr::mutate(decreasing = c(FALSE, diff(.data$density) < 0),
    localextremum = c(FALSE, diff(.data$decreasing) != 0),
    nblocalextrema = cumsum(.data$localextremum)) |>
  dplyr::filter(.data$decreasing) |>
  dplyr::group_by(.data$nblocalextrema) |>
  dplyr::slice(1) |>
  dplyr::arrange(desc(density)) |>
  dplyr::pull(.data$abciss) |>
  list()
}
```


Explore

This is a fairly small dataset. Variables of interest are `artists_share`, the percentage of workers that are artists, and `location_quotient`, the ratio between a certain state's number of workers and the overall US's. `location_quotient` can be roughly thought of as where an occupation gravitates to. We are also interested in the distribution of values across the US.


To start, Designers are the most popular artist by a long shot, in terms of absolute numbers:

```
df %>% group_by(type) %>% summarize(sum(artists_n)) %>% arrange(desc(`sum(artists_n)`))
```

```
# A tibble: 13 × 2
  type                                `sum(artists_n)`
  <chr>                                <dbl>
1 Designers                          935090
2 Writers And Authors                252870
3 Fine Artists, Art Directors, And Animators 235915
4 Photographers                      191120
5 Architects                         185495
6 Producers And Directors            179995
7 Musicians                         170949
8 Announcers                        81850
9 Actors                            68255
10 Entertainers                      62115
11 Music Directors And Composers       59320
12 Landscape Architects               42579
13 Dancers And Choreographers         30220
```

We can see that there is not much difference in `artists_share` across both race and region, note the square root x scale.

```
df %>%
  ggplot(aes(x = artists_share, color = race)) +
  geom_boxplot() +
  scale_x_sqrt(breaks = c(1/1000, 1/500, 1/250, 1/125, 1/62.5),
    labels = scales::label_percent(0.01)) + facet_wrap(vars(region)) +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1))
```




We can also inspect how the occupation types differ in relative popularity, the `artists_share`. The plots are split by rough order, for clarity. Here, we can see just how much many more Designers there are, even compared to the next closest group, Fine Artists, Art Directors, and Animators.

```

modes <- df %>%
  select(type, artists_share) %>%
  group_by(type) %>%
  summarize(mode = ModEstM(artists_share) %>% first %>% first) %>%
  mutate(mode_lvls = cut_number(mode, 4, labels = c("Low", "Med-Low", "Med-High",
    ↪ "High"))))

df %>%
  left_join(modes, by="type") %>%
  ggplot(aes(x = artists_share)) +
  geom_density(aes(color = type)) +
  scale_x_continuous(trans = "sqrt", limits = c(0, .015),
    breaks = c(1/10000, 1/2000, 1/1000, 1/500, 1/250, 1/125,
    ↪ 1/62.5),
    labels = scales::label_percent(0.01)) +
    facet_wrap(vars(mode_lvls), ncol = 1) +
  theme(legend.text = element_text(size=6), legend.key.size = unit(4, 'mm'))

```



GIS

Create a dataframe with states as rows, for use later in GIS.

```
df_state <- df %>% group_by(state) %>%
  summarize(share = mean(artists_share),
            quot = mean(location_quotient),
            work = sum(all_workers_n),
            arts = sum(artists_n))

df_state %>%
  arrange(desc(quot))
```

A tibble: 52 × 5

	state	share	quot	work	arts
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	District of Columbia	0.00366	2.06	5302765	14195
2	California	0.00173	1.83	258980085	445095
3	Nevada	0.00128	1.72	19596005	25440
4	New York	0.00170	1.72	130899860	243345
5	Hawaii	0.00153	1.32	9776845	13970
6	Maryland	0.00108	1.17	42499990	48030
7	Oregon	0.00211	1.15	27256970	42350
8	Georgia	0.00110	1.14	67268955	67940
9	Colorado	0.00104	1.12	39905255	51410
10	Florida	0.000922	1.11	131508390	139505

... with 42 more rows

Import map data and combine with original data.

```
library(rnaturalearthhires)

usa <- rnaturalearth::ne_states("United States of America", returnclass = "sf")
usa_df <- right_join(usa, df_state,
                    by = c("name" = "state")) %>%
  rename(`Location Quotient` = quot,
        `# All Workers` = work,
        `# Artists` = arts,
        `Artist Proportion` = share)
```

Map out the four relevant variables:

```

## whole us: xlim = c(-180, -65)

## make function over fill and palette
## fill = c(share, quot, work, arts)
constructMap <- function(fill = quot){
  palette <- switch(fill,
    "Location Quotient" = scale_fill_gradient2(high = "#b2182b", #
      ↪ colorbrewer RdGy
      mid = "#ffffff",
      low = "#00ffff", midpoint
      ↪ = 1),
    "# All Workers" = scale_fill_distiller(palette = "Reds", direction =
      ↪ 1,
      label = scales::label_number(suffix =
        ↪ "M", scale = 1e-6)),
    "# Artists" = scale_fill_distiller(palette = "Blues", direction = 1,
      label = scales::label_number(suffix =
        ↪ "K", scale = 1e-3)),
    "Artist Proportion" = scale_fill_distiller(palette = "Purples",
      ↪ direction = 1,
      labels = scales::label_percent(0.01)),
    stop("Not a valid fill value"))

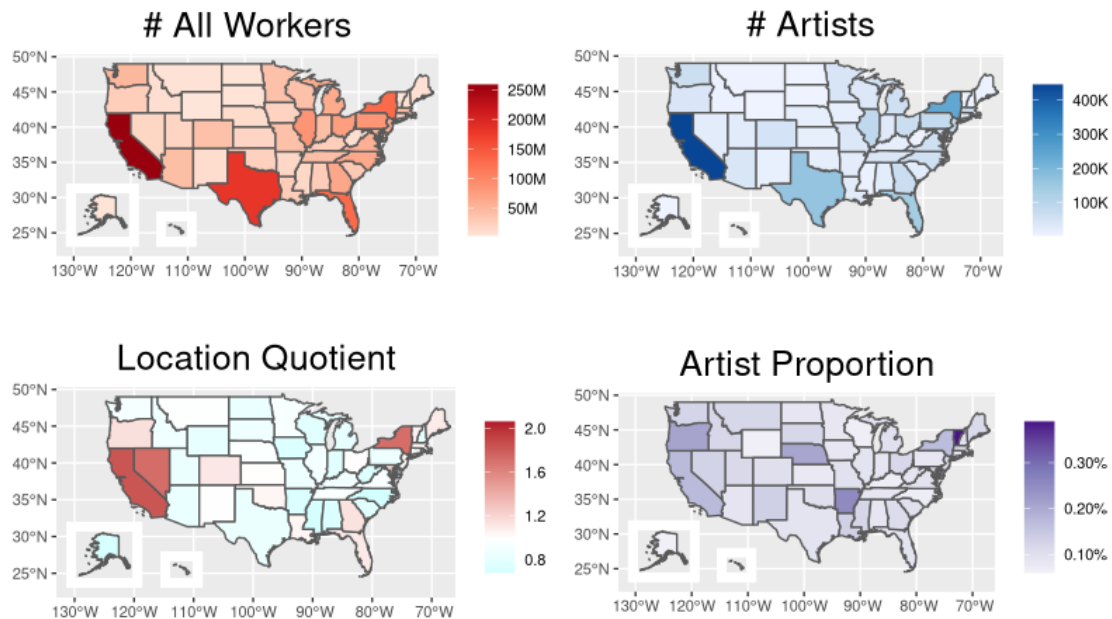
world <- usa_df %>%
  ggplot(aes(fill = !!sym(fill))) +
  geom_sf() +
  palette

mainland <- world + coord_sf(xlim = c(-130, -69), ylim = c(23, 49)) +
  ggtitle(fill) + theme(plot.title = element_text(size = 20, hjust = 0.5),
    ↪ legend.title=element_blank())
alaska <- world + coord_sf(xlim = c(-180, -130), ylim = c(51, 71), datum = NA) +
  ↪ theme(legend.position="none")
hawaii <- world + coord_sf(xlim = c(-161, -154), ylim = c(18, 23), datum = NA) +
  ↪ theme(legend.position="none")

mainland +
  annotation_custom(
    grob = ggplotGrob(alaska),
    xmin = -140, xmax = -110,
    ymin = 23, ymax = 32) +
  annotation_custom(
    grob = ggplotGrob(hawaii),
    xmin = -124, xmax = -100,
    ymin = 23, ymax = 28)
}

gridExtra::grid.arrange(constructMap("# All Workers"), constructMap("# Artists"),
  constructMap("Location Quotient"), constructMap("Artist
  ↪ Proportion"))

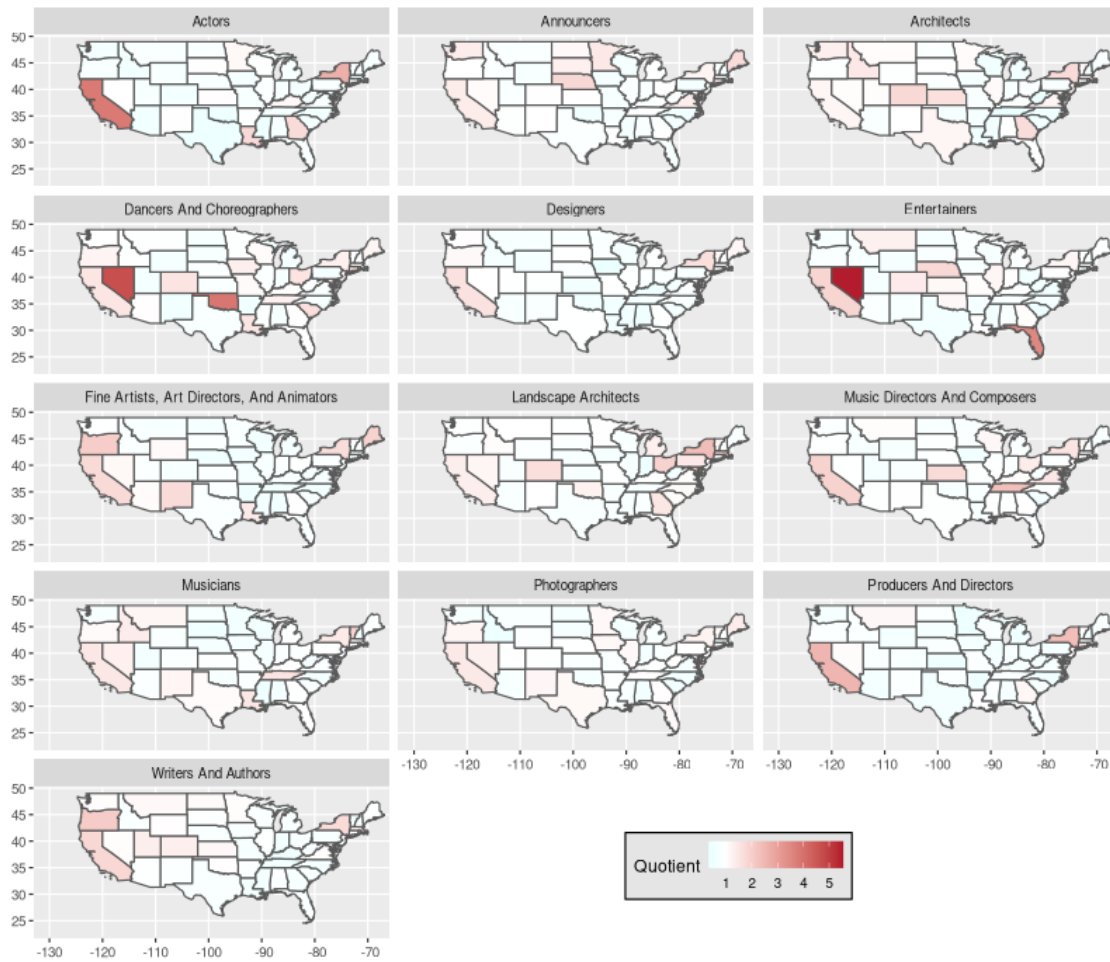
```



We can see that both the total number of workers and total number of artists distribute similarly, primarily in California, Texas, New York, and Florida. Likewise, this means the proportion of artists is fairly uniform across the country. The location quotient tells us that artists in general concentrate in California, Nevada, and New York.

To inspect this quotient relationship, plot by each artist type:

```
df %>%
  group_by(state, type) %>%
  summarize(Quotient = mean(location_quotient)) %>%
  rename(job = type) %>%
  right_join(usa, by = c("state" = "name")) %>%
  ggplot(aes(fill = Quotient)) +
  geom_sf(aes(geometry = geometry)) +
  facet_wrap(vars(job), ncol = 3) +
  scale_fill_gradient2(high = "#b2182b", # colorbrewer RdGy
                       mid = "#ffffff",
                       low = "#00ffff", midpoint = 1) +
  coord_sf(xlim = c(-130, -69), ylim = c(23, 49)) +
  theme(legend.position = c(0.65, 0.08),
        legend.direction = "horizontal",
        legend.background = element_rect(fill = "gray90", colour = "black"))
```



Some interesting points:

- Actors & Producers And Directors congregate in California, New York, and Georgia
- Nevada is where Dancers and Choreographers & Entertainers are
- The scale is highly skewed by Actors, Dancers And Choreographers, & Entertainers, which I attribute to Hollywood and Las Vegas.

Resources

- MICE: https://www.gerkovink.com/miceVignettes/Convergence_pooling/Convergence_and_pooling.html
- <https://r-spatial.org/r/2018/10/25/ggplot2-sf-3.html>